

# Forecasting Lifetime Value (LTV) using PySTAN



# About me

- **2022 ~ Today: Data Science Lead in Marketing Science at Meta**
  - Apply causal inference techniques to extract insight to help advertisers in increasing the effectiveness of their marketing investments
  - Maintainer of GeoLift, an open-source techniques for measuring marketing effectiveness
- **2017 ~ 2022: Staff Data Scientist in Wildlife Studios**
  - Tech lead for the Lifetime Value team
  - Developed the Lifetime Value models for launching 5 games
  - Implemented ads monetization algorithm which increases daily ads revenue



# Topics


- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - The data used: Lifetime Value dataset from Kaggle
  - Description of the model
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC




# Topics

- Introduction to Advertisement
  - **Modelling advertisement for digital products**
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - The data used: Lifetime Value dataset from Kaggle
  - Description of the model
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC







# Day 1




## Day 2




Day ...












Unreached





Organic
































Day 1






Day 1



Day 2










Day 1


Day 2

Day 3






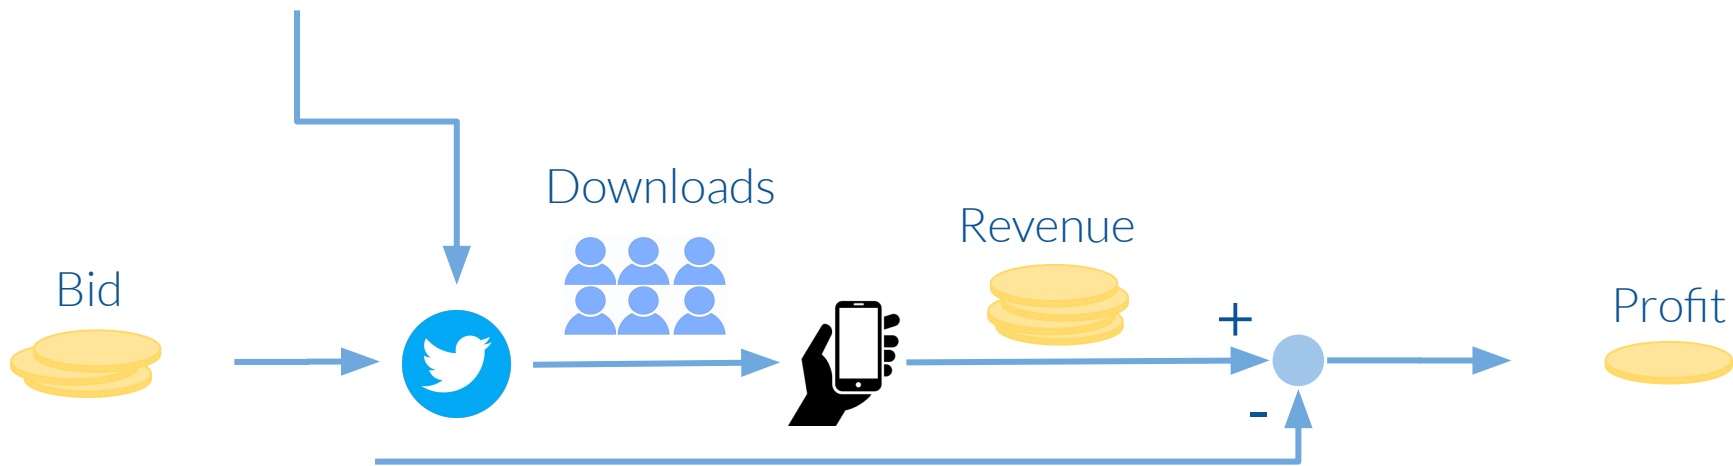











Audience a.k.a.  
Targeted Population







- Volume Function: How does the bid affects number of downloads?



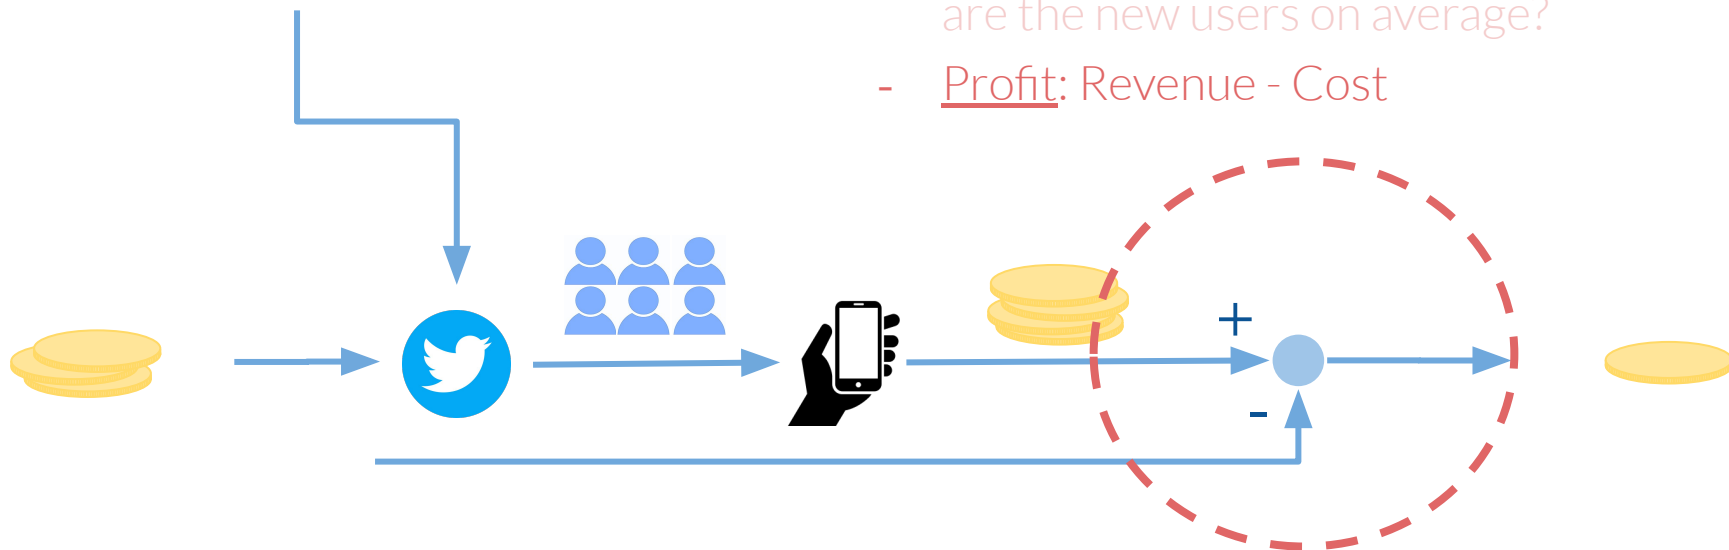



- Volume Function: How does the bid affects number of downloads?
- Lifetime Value (LTV): How much worth are the new users on average?



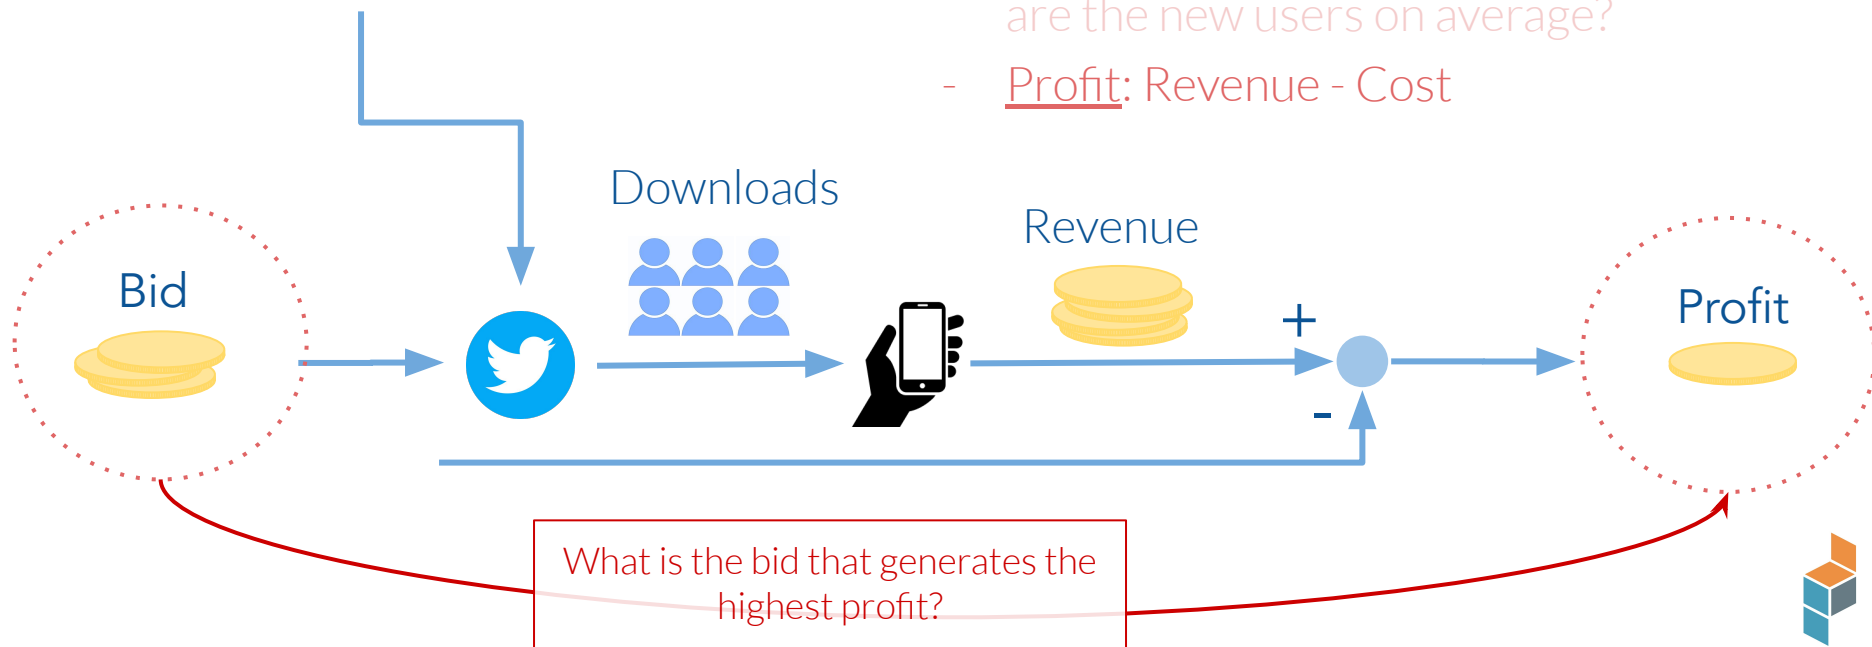


- Volume Function: How does the bid affects number of downloads?
- Lifetime Value (LTV): How much worth are the new users on average?
- Profit: Revenue - Cost






- Volume Function: How does the bid affects number of downloads?
- Lifetime Value (LTV): How much worth are the new users on average?
- Profit:  $\text{Revenue} - \text{Cost}$



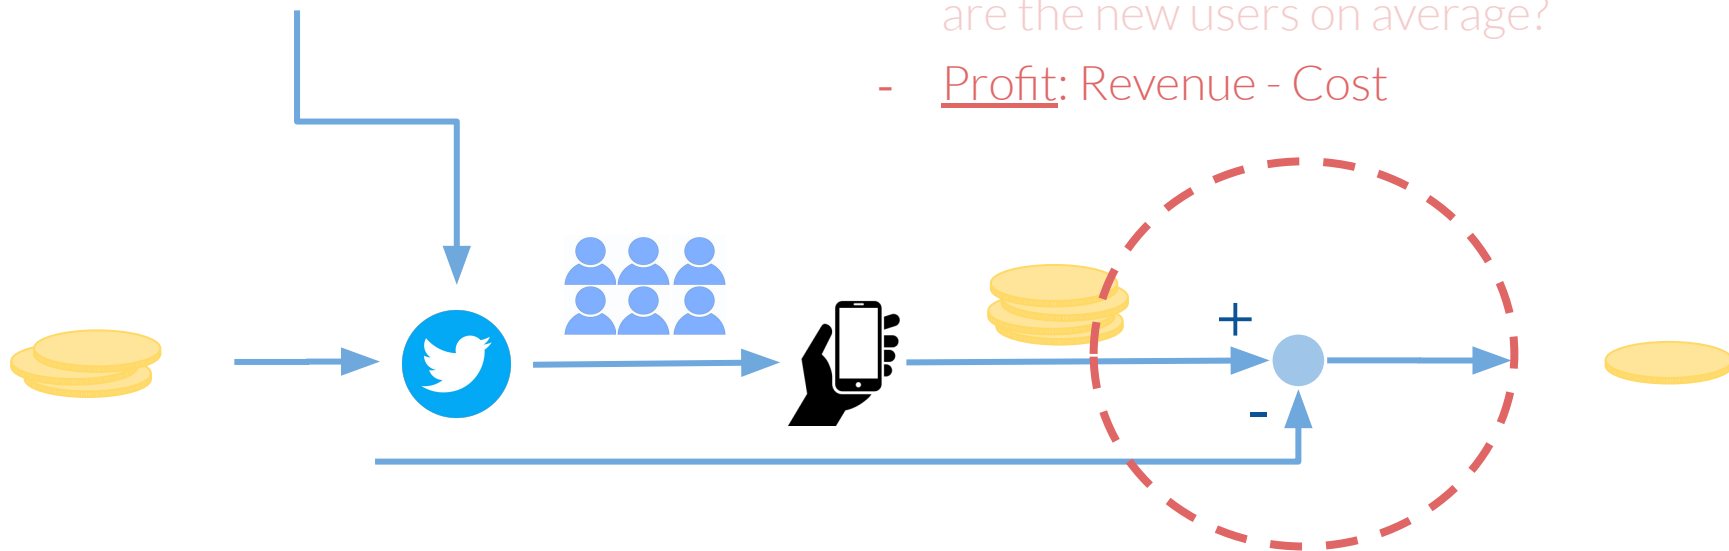
# Topics

- Introduction to Advertisement
  - Modelling advertisement for digital products
  - **Finding the optimal bid for your marketing campaigns**
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - The data used: Lifetime Value dataset from Kaggle
  - Description of the model
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC





- Volume Function: How does the bid affects number of downloads?
- Lifetime Value (LTV): How much worth are the new users on average?
- Profit: Revenue - Cost



$$\textit{Profit} = \textit{Revenue} - \textit{Cost}$$





$$Profit = Revenue - Cost$$

$$Profit = LTV * Downloads - Bid * Downloads$$



$$Profit = Revenue - Cost$$

$$Profit = LTV * Downloads - Bid * Downloads$$

$$Profit = (LTV - Bid) * Downloads$$



$$Profit = Revenue - Cost$$

$$Profit = LTV * Downloads - Bid * Downloads$$

$$Profit = (LTV - Bid) * Downloads$$

$$Profit = (LTV - Bid) * Vol(Bid)$$



$$Profit = Revenue - Cost$$


$$Profit = LTV * Downloads - Bid * Downloads$$


$$Profit = (LTV - Bid) * Downloads$$


$$Profit = (LTV - Bid) * Vol(Bid)$$


$$Profit = f(Bid)$$



















# Topics


- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - **The role that uncertainty play in the marketing strategy**
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - The data used: Lifetime Value dataset from Kaggle
  - Description of the model
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC



















% from max profit

100%

0%

-100%

-200%

-300%

Error



% from max profit

Error

When the Profit Curve is asymmetric around its optimal value - usually a consequence of the Volume Function - you may need to underbid on your marketing campaigns



% from max profit

100%

small error

0%

-100%

large error

-200%

-300%

Error

When the we are uncertain about the expected LTV, we also need to be careful on how we bid



# Topics

- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - **What is PySTAN**
  - The data used: Lifetime Value dataset from Kaggle
  - Description of the model
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC



# What is (Py)STAN?

PySTAN is a Python interface to STAN.

STAN itself is a platform statistical modelling and allows high-performance statistical computation. It was first released in 2012, and since then has seen wide use in different scientific fields because of its interface to both R and Python, the implementation of efficient sampling algorithms, and usage of automatic differentiation



# Topics

- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - **The data used: Lifetime Value dataset from Kaggle**
  - Description of the model
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC



≡ kaggle

+ Create

 Home

## Competitions

## Datasets

 Models

Code

Discussions

 Learn

More



**Sign In**

## Register



DAVID RODRÍGUEZ SEGADO · UPDATED 3 YEARS AGO

▲ 9

## New Notebook

Download (11 MB)



## A decorative graphic consisting of several overlapping geometric shapes in bright colors: red, orange, yellow, green, and blue. The shapes include triangles, squares, and hexagons, creating a vibrant, abstract pattern.

Data Card	Code (2)	Discussion (0)
-----------	----------	----------------

**Explanation of the data set:**

The files contain randomized data about the behaviour and spending patterns of customers. This data is used to model the amount a user will spend during his membership on specific websites. This metric is called customer lifetime value (LTV). The training set data are from users that have joined our website from Dec/2018 through June/2019. The test set contains users that joined in July/2019. The data set contains the following columns:

- **product\_type**: one of 4 different product\_type a user can purchase. The user always gets a first product (called 'type\_p') and then can purchase additional products (of 'type\_x', 'type\_u', and 'type\_ex').
- **user\_id**: a numerical identifier of the user. Note that for every user we might

**Usability** ⓘ

5.59

## License

Unknown

**Expected update frequency**

Not specified

## Tags

## Business

# The dataset used: Lifetime Value from Kaggle

	product_type	user_id	join_date	hidden	product	STV	target	credit_card_level	is_lp	aff_type	is_cancelled	country_segment
0	type_ex	7.0	2018-12-01 00:01:45	0	product_1	8.25	8.25	standard	0	PPL	NaN	US
1	type_ex	20.0	2018-12-01 00:06:05	0	product_2	8.25	8.25	standard	0	PPL	NaN	US
2	type_ex	22.0	2018-12-01 00:06:23	0	product_3	8.25	8.25	prepaid	0	PPL	NaN	US
3	type_ex	26.0	2018-12-01 00:07:12	0	product_2	8.25	8.25	standard	0	PPL	NaN	US
4	type_ex	59.0	2018-12-01 00:15:21	0	product_2	8.25	8.25	standard	0	PPL	NaN	Other Countries
5	type_ex	63.0	2018-12-01 00:15:28	0	product_1	8.25	8.25	standard	0	PPS	NaN	US
6	type_ex	70.0	2018-12-01 00:16:32	0	product_4	8.25	8.25	standard	0	PPL	NaN	US
7	type_ex	82.0	2018-12-01 00:21:14	0	product_1	8.25	8.25	standard	0	PPS	NaN	US
8	type_ex	87.0	2018-12-01 00:22:40	0	product_2	8.25	8.25	standard	1	PPS	NaN	Other Countries
9	type_ex	102.0	2018-12-01 00:26:55	0	product_2	8.25	8.25	standard	0	PPL	NaN	US






# The dataset used: Lifetime Value from Kaggle


	product_type	user_id	join_date	hidden	product	STV	target	credit_card_level	is_lp	aff_type	is_cancelled	country_segment
0	type_ex	7.0	2018-12-01 00:01:45	0	product_1	8.25	8.25	standard	0	PPL	NaN	US
1	type_ex	20.0	2018-12-01 00:06:05	0	product_2	8.25	8.25	standard	0	PPL	NaN	US
2	type_ex	22.0	2018-12-01 00:06:23	0	product_3	8.25	8.25	prepaid	0	PPL	NaN	US
3	type_ex	26.0	2018-12-01 00:07:12	0	product_2	8.25	8.25	standard	0	PPL	NaN	US
4	type_ex	59.0	2018-12-01 00:15:21	0	product_2	8.25	8.25	standard	0	PPL	NaN	Other Countries
5	type_ex	63.0	2018-12-01 00:15:28	0	product_1	8.25	8.25	standard	0	PPS	NaN	US
6	type_ex	70.0	2018-12-01 00:16:32	0	product_4	8.25	8.25	standard	0	PPL	NaN	US
7	type_ex	82.0	2018-12-01 00:21:14	0	product_1	8.25	8.25	standard	0	PPS	NaN	US
8	type_ex	87.0	2018-12-01 00:22:40	0	product_2	8.25	8.25	standard	1	PPS	NaN	Other Countries
9	type_ex	102.0	2018-12-01 00:26:55	0	product_2	8.25	8.25	standard	0	PPL	NaN	US



# The dataset used: Lifetime Value from Kaggle




# The dataset used: Lifetime Value from Kaggle



# Topics

- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - The data used: Lifetime Value dataset from Kaggle
  - **Description of the model**
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC






Day 1




Day 2



Day 3





Day 1




Day 2



Day 3



Day 1




Day 2





Day 3



Day 1



Day 2





Day 3









Day 1




Day 2




Day 3





Day 1




Day 2



Day 3




$$LTV_{Population}(t) \sim Normal \left( LTV_{Population}(t-1), \sigma^2 \right)$$



**Random  
Walk**



$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$



# The (Py)STAN Model

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

Modelling population LTV as a Random Walk

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$

LTV of a user is a 'sample' of the population LTV

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

Prior for population LTV



# Topics

- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - The data used: Lifetime Value dataset from Kaggle
  - Description of the model
  - **Implementation of the model in PySTAN**
  - Achieving the same model with PyMC
  - Comparison between PySTAN and PyMC



# The (Py)STAN Model

```
import stan

model = stan.build(time_series_model, data=stan_data)
posterior_distributions = model.sample(numchains=3, num_samples=1000)
```

*Note: In the original image, 'time\_series\_model' is highlighted with a red dashed box labeled 'str' and 'stan\_data' is highlighted with a blue dashed box labeled 'dict'.*



# The (Py)STAN Model

```
import stan

model = stan.build(time_series_model, data=stan_data)
posterior_distributions = model.sample(numchains=3, num_samples=1000)
```

```
data = pd.read_csv('lifetime_value.csv')

stan_data = {}
stan_data['N'] = data.shape[0]
stan_data['n_dates'] = len(data['join_date'].unique())
stan_data['date'] = data['join_date']
stan_data['observation'] = data['target']
```





# The (Py)STAN Model



```
import stan

model = stan.build(time_series_model, data=stan_data)
posterior_distributions = model.sample(numchains=3, num_samples=1000)
```



# The (Py)STAN Model

```
time_series_model = """
    data {

    }

    parameters {

    }

    model {

    }
    """
```



# The (Py)STAN Model

```
time_series_model = """
    data {

    }

    parameters {

    }

    model {

    }
    """
```

**Data** defines what are the data and constants passed to the model

**Parameters** declare the variables you want to estimate

**Model** defines the connection between the parameters and the data



# The (Py)STAN Model

```
time_series_model = """
    data {
        int<lower=0> N;
        int<lower=0> n_dates;
        int date[N];
        real observation[N];
    }
    parameters {

    }
    model {

    }
}
"""
```

```
data = pd.read_csv('lifetime_value.csv')

stan_data = {}
stan_data['N'] = data.shape[0]
stan_data['n_dates'] = len(data['join_date'].unique())
stan_data['date'] = data['join_date']
stan_data['observation'] = data['target']
```



# The (Py)STAN Model

```
time_series_model = """
    data {
        int<lower=0> N;
        int<lower=0> n_dates;
        int date[N];
        real observation[N];
    }
    parameters {

    }
    model {

    }
}
"""
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$



# The (Py)STAN Model

```
time_series_model = """
  data {
    int<lower=0> N;
    int<lower=0> n_dates;
    int date[N];
    real observation[N];
  }
  parameters {
    vector[n_dates] mu;

  }
  model {

  }
  """
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$



# The (Py)STAN Model

```
time_series_model = """
  data {
    int<lower=0> N;
    int<lower=0> n_dates;
    int date[N];
    real observation[N];
  }
  parameters {
    vector[n_dates] mu;
    real<lower=0.0001> sampling_stddev;

  }
  model {

  }
  """
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$

# The (Py)STAN Model

```
time_series_model = """
  data {
    int<lower=0> N;
    int<lower=0> n_dates;
    int date[N];
    real observation[N];
  }
  parameters {
    vector[n_dates] mu;
    real<lower=0.0001> sampling_stddev;
    real<lower=0.0001> random_walk_stddev;
  }
  model {

  }
}
"""
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$





# The (Py)STAN Model

```
time_series_model = """
  data {
    int<lower=0> N;
    int<lower=0> n_dates;
    int date[N];
    real observation[N];
  }
  parameters {
    vector[n_dates] mu;
    real<lower=0.0001> sampling_stddev;
    real<lower=0.0001> random_walk_stddev;
  }
  model {
    sampling_stddev ~ cauchy(0, 2);
    random_walk_stddev ~ cauchy(0, 2);
    mu[1] ~ normal(6, 3);
  }
  """
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$



# The (Py)STAN Model

```
time_series_model = """
  data {
    int<lower=0> N;
    int<lower=0> n_dates;
    int date[N];
    real observation[N];
  }
  parameters {
    vector[n_dates] mu;
    real<lower=0.0001> sampling_stddev;
    real<lower=0.0001> random_walk_stddev;
  }
  model {
    sampling_stddev ~ cauchy(0, 2);
    random_walk_stddev ~ cauchy(0, 2);
    mu[1] ~ normal(6, 3);
    mu[2:n_dates] ~ normal(mu[1:(n_dates - 1)], random_walk_stddev);
  }
  """
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t - 1), \sigma_{Time}^2)$$

$$LTV_{Population}(t = 0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$



# The (Py)STAN Model

```
time_series_model = """
  data {
    int<lower=0> N;
    int<lower=0> n_dates;
    int date[N];
    real observation[N];
  }
  parameters {
    vector[n_dates] mu;
    real<lower=0.0001> sampling_stddev;
    real<lower=0.0001> random_walk_stddev;
  }
  model {
    sampling_stddev ~ cauchy(0, 2);
    random_walk_stddev ~ cauchy(0, 2);
    mu[1] ~ normal(6, 3);
    mu[2:n_dates] ~ normal(mu[1:(n_dates - 1)], random_walk_stddev);

    observation ~ normal(mu[date], sampling_stddev);
  }
  """
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t - 1), \sigma_{Time}^2)$$

$$LTV_{Population}(t = 0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$




# The (Py)STAN Model

```
import stan


model = stan.build(time_series_model, data=stan_data)
posterior_distributions = model.sample(numchains=3, num_samples=1000)
```




# The (Py)STAN Model




# The (Py)STAN Model



# The (Py)STAN Model




# The (Py)STAN Model





# The (Py)STAN Model



# Topics

- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - Description of the model
  - The data used: Lifetime Value dataset from Kaggle
  - Implementation of the model in PySTAN
  - **Achieving the same model with PyMC**
  - Comparison between PySTAN and PyMC



# The PyMC Model

```
import pymc as pm

coords = {"steps": data['join_date'].values}
with pm.Model(coords=coords) as model:
```



# The PyMC Model

```
import pymc as pm
```

```
coords = {"steps": data['join_date'].values}
```

```
with pm.Model(coords=coords) as model:
```

```
# Priors on Gaussian random walks
```

```
mu_prior = pm.Normal.dist(6, 3)
```

```
random_walk_stddev = pm.HalfCauchy('random_walk_stddev', 2)
```

```
sampling_stddev = pm.HalfCauchy("sampling_stddev", 2)
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$



# The PyMC Model

```
import pymc as pm
```

```
coords = {"steps": data['join_date'].values}
with pm.Model(coords=coords) as model:
```

```
    # Priors on Gaussian random walks
```

```
    mu_prior = pm.Normal.dist(6, 3)
```

```
    random_walk_stddev = pm.HalfCauchy('random_walk_stddev', 2)
```

```
    sampling_stddev = pm.HalfCauchy("sampling_stddev", 2)
```

```
    alpha = pm.Deterministic(
        "alpha", pt.concatenate([alpha_prior,
        pm.Normal("alpha_raw", sigma=alpha_dev, shape=n_days-1)]).cumsum(axis=0)
    )
```

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$



# The PyMC Model

$$LTV_{Population}(t) \sim Normal(LTV_{Population}(t-1), \sigma_{Time}^2)$$

$$LTV_{Population}(t=0) \sim Normal(\mu_{Prior}, \sigma_{Prior}^2)$$

$$LTV_{User}(t) \sim Normal(LTV_{Population}(t), \sigma_{Sampling}^2)$$

```
import pymc as pm
```

```
coords = {"steps": data['join_date'].values}
```

```
with pm.Model(coords=coords) as model:
```

```
# Priors on Gaussian random walks
```

```
mu_prior = pm.Normal.dist(6, 3)
```

```
random_walk_stddev = pm.HalfCauchy('random_walk_stddev', 2)
```

```
sampling_stddev = pm.HalfCauchy("sampling_stddev", 2)
```

```
alpha = pm.Deterministic(
```

```
    "alpha", pt.concatenate([alpha_prior,
```

```
    pm.Normal("alpha_raw", sigma=alpha_dev, shape=n_days-1)]).cumsum(axis=0)
```

```
)
```

```
likelihood = pm.Normal(
```

```
    "likelihood", mu=mu[dates_tensor], sigma=sampling_stddev, observed=target_tensor, dims="steps"
```

```
)
```



# Topics

- Introduction to Advertisement
  - Modelling advertisement for digital products
  - Finding the optimal bid for your marketing campaigns
  - The role that uncertainty play in the marketing strategy
- Forecasting Lifetime Value with PySTAN
  - What is PySTAN
  - The data used: Lifetime Value dataset from Kaggle
  - Description of the model
  - Implementation of the model in PySTAN
  - Achieving the same model with PyMC
  - **Comparison between PySTAN and PyMC**



# Comparison between PySTAN and PyMC

Using the Lifetime Value dataset for the tests, PySTAN outperforms PyMC for 'small' datasets and low-cardinality categories when using the **default samplers**. This is coherent with a benchmark conducted by [Joshua Cook](#) comparing STAN and PyMC3\*


- As the data becomes larger (>20k) PyMC starts to performs better than PySTAN

\* Joshua Cook's benchmark compares PySTAN vs PyMC3, which is an older version of PyMC to what we have today and less optimized





# Comparison between PySTAN and PyMC



# Comparison between PySTAN and PyMC


Using the Lifetime Value dataset for the tests, PySTAN outperforms PyMC for 'small' datasets and low-cardinality categories when using the **default samplers**. This is coherent with a benchmark conducted by [Joshua Cook](#) comparing STAN and PyMC3\*

- As the data becomes larger (>20k) PyMC starts to performs better than PySTAN
- If using '*numpyro*' or '*nutpie*' samplers from PyMC, performance jumps up to 3x

\* Joshua Cook's benchmark compares PySTAN vs PyMC3, which is an older version of PyMC to what we have today and less optimized



# Comparison between PySTAN and PyMC



# Comparison between PySTAN and PyMC

One point worth of attention is that PyMC easily allows for multiple implementation of the same model through different classes. However, some of this classes can display significantly lower performance.



# Comparison between PySTAN and PyMC


```
alpha = pm.Deterministic(  
    "alpha", pt.concatenate([alpha_prior,  
        pm.Normal("alpha_raw", sigma=alpha_dev, shape=n_days-1)]).cumsum(axis=0)  
)
```

VS


```
alpha = pm.GaussianRandomWalk(  
    "alpha", mu=0, sigma=alpha_dev, init_dist=pm.Normal.dist(6, 3), shape=n_days  
)
```



# Comparison between PySTAN and PyMC



# Comparison between PySTAN and PyMC(3)



\*Benchmark by Joshua Cook



# Comparison between PySTAN and PyMC

In summary

- Use PyMC if
  - Your dataset is large ( >20K data points),
  - Models are complex (e.g. hierarchical models, high cardinality, many covariates),
  - There is a ready to use implementation available (e.g. PyMC Marketing Mix Models)
  - You can use in-development samplers numpyro or nutpie (require extra installations)
- Use PySTAN if
  - Your dataset isn't too large
  - Your model isn't too complex
  - You or colleges have experience working with STAN from another programming library (e.g. R)






# Conclusions

- Depending on the properties of your marketing campaigns, the uncertainty on your estimates change the strategy in operating them
- We saw how we can use PyMC or PySTAN to forecast LTV
- PyMC performs better than PySTAN for larger datasets or more complex models
- PyMC can have a steeper learning curve, but can be worthwhile if time is not critical



Special thanks



Thomas  
Wiecki



PyMC  
Labs

THE BAYESIAN CONSULTANCY

[www.pymc-labs.io](http://www.pymc-labs.io)




Q & A





# PyMC Faster Sampling with JAX and Numba





# Comparison between PySTAN and PyMC - Speed

From my experience, and from benchmark conducted by [Joshua Cook](#), PySTAN still outperforms PyMC for 'small' datasets and low-cardinality categories



[https://github.com/pymc-devs/pymc-examples/blob/23e42c055c8203f1a57dd5726033b6068d02718b/examples/samplers/fast\\_sampling\\_with\\_jax\\_and\\_numba.ipynb](https://github.com/pymc-devs/pymc-examples/blob/23e42c055c8203f1a57dd5726033b6068d02718b/examples/samplers/fast_sampling_with_jax_and_numba.ipynb)



# Comparison between PySTAN and PyMC - Resources

A perfect comparison on the number and quality of available resources on PySTAN and PyMC is never going to be perfect, but we can use some reference points:

- Average daily downloads ([PyPiStats](#)):
  - PySTAN 57.8k
  - PyMC: 19.2k
- Github stars
  - PySTAN: 0.3k
  - PyMC: 7.7k
- Ready to use models
  - PySTAN: Only statistical distributions
  - PyMC: many ready to use models (e.g. LTV, Medix Mix Modelling)





# The (Py)STAN Model - Adapting the data for PySTAN

```
stan_format_data = {}  
stan_format_data['N'] = int(data.shape[0])  
stan_format_data['n_dates'] = len(data['join_date'].unique())  
stan_format_data['date'] = list(data['join_date'].values)  
stan_format_data['observation'] = list(data['target'].values)
```

PySTAN receives a  
data only as a  
dictionary



# The PyMC Model

```
dates_tensor = pytensor.shared(data['join_date'].values)
target_tensor = pytensor.shared(data['target'].values)
n_days = len(data['join_date'].unique())
```



Similarly, PyMC receives data as 'independent' vectors




# How to adapt the bidding strategy depending on the uncertainty around LTV

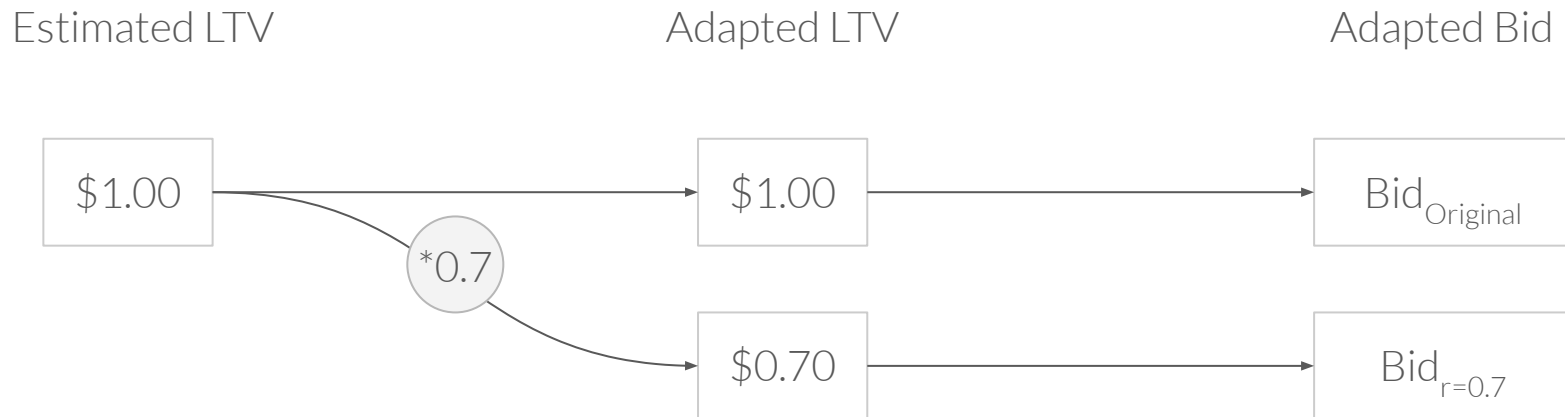
Estimated LTV

Adapted LTV

Adapted Bid



# How to adapt the bidding strategy depending on the uncertainty around LTV




# How to adapt the bidding strategy depending on the uncertainty around LTV

Estimated LTV

Adapted LTV

Adapted Bid



Fraction of optimal bidding point that provides largest profit

1.0

0.8

0.6

0.4

0.2

Symmetric Profit Curve and low-sensibility Volume Function

Asymmetric Profit Curve and high-sensibility Volume Function

Error



Fraction of optimal bidding point that provides largest profit

1.0

0.8

0.6

0.4

0.2

Symmetric Profit Curve and low-sensibility Volume Function

Asymmetric Profit Curve and high-sensibility Volume Function

Error



Fraction of optimal bidding point that provides largest profit

1.0

0.8

0.6

0.4

0.2

Symmetric Profit Curve and low-sensibility Volume Function

Asymmetric Profit Curve and high-sensibility Volume Function

Error

